Leak or Keep?
Learning to Share Interview Questions Among Competing Candidates

Final Project Report — Multi-Agent Systems, IIT Gandhinagar

Jash Bharat Pasad 24110146 jash.pasad@iitgn.ac.in

Bhavay Goyal 24110070 bhavay.goyal@iitgn.ac.in

Chaitanya Bhoite 24110086 chaitanya.bhoite@iitgn.ac.in

Vishnu Sinha 24110389 vishnu.sinha@iitgn.ac.in

Code, result files and commands that regenerate every number below:
https://github.com/Vish-Overflows/multiagent-systems-leak-or-keep

Abstract. After an interview, a candidate can tell friends who interview later which questions were asked. This helps them, but they may be competing for the same offers. We model this as a game on an $n \times n$ sharing matrix: candidates interview in a fixed order, know only a coarse ability type (never their score), and a shared question may not reappear. We solve the six-candidate game *exactly* for 378 combinations of hiring rule, friendliness, friendship network and leak model, and then let bandit learners who know none of the rules play it for 200,000 rounds (1080 runs). Four findings. (1) *The hiring rule decides who talks.* If the company hires everyone above a cut-off, everybody tells as soon as they care about friends at all. If it fills a fixed number of seats, only candidates with little to lose tell: weak ones first (once they value a friend's offer at 0.10 of their own), strong ones last (0.38), and a candidate who is friends with everyone never. (2) Disliking outsiders makes candidates tell *more* under fixed seats and changes nothing under a cut-off, as our theory predicts. (3) Learners who see only their own outcome find this equilibrium: UCB1 ends with 93–99% of decisions correct, and learners that use how leaks work (a tell that did not land is the same as no tell; or an estimated leak rate and value per leaked question) get there two to five times sooner. (4) At friendliness 0.2, leaks cost a cut-off company 8–12 times more hiring value than a fixed-seats company: it needs a bank of 40 questions, against 9, to keep the loss under 5%. Most of that gap is *who* tells (everyone under a cut-off, only weak candidates under fixed seats): if everyone told under both rules, the cut-off would lose 2.5–4.4 times more and need 40 questions against 20.

Keywords Bayesian games · multi-agent bandits · empirical game theory · coarse correlated equilibrium

1 Introduction

Every placement season, a candidate leaves the interview room and within minutes their friends know what was asked. It feels like helping friends, but when offers are limited, a friend helped may take the helper's seat. Leaks happen at scale [1, 2]; industry treats them as something that happens to questions (bigger banks, rotation [2]), not as a choice made by competing candidates.

Example. Two offers; Aarav, Meera, Kabir and Diya interview in that order, and Kabir and Diya are Aarav's friends. Each answers 3 of 10 bank questions, and knowing a question adds 4 of its 10 marks. With scores 21, 22, 20 and 19, Meera and Aarav are hired. If Aarav shares, each friend sees at least one of his questions with probability $1 - \binom{7}{3}/\binom{10}{3} \approx 0.71$, and with probability 0.92 one of them overtakes him. If the company instead hires everyone scoring at least 20, sharing costs Aarav nothing — and Diya, meant to be rejected, gets in with probability 0.71.

Question. Who shares with whom, how does this depend on the hiring rule and on what candidates care about, and can candidates who do not know the payoffs *learn* to share as the equilibrium predicts?

Related work. Rank-order versus absolute rewards [3] and helping or sabotaging rivals [4] are well studied, but there help is effort, not information. Li and Conitzer [5] model leaked test questions as a Stackelberg game with the leak given; we make it a choice among competitors. With unknown payoffs, candidates are bandit learners [6, 7, 8], whose no-regret play approaches coarse correlated equilibria [9]; solving a simulated game is empirical game theory [10].

What this report adds to the project review. The review fixed the model, proved Lemma 1 and Propositions 1–3, and showed a one-setting pilot. Here we describe the algorithms (Section 4), report all four planned experiments (Section 5), and give a verdict on each of the five hypotheses fixed in advance (Table 4). We use two plain names: *fixed seats* (Seats: the best K scores are hired) and *cut-off* (Bar: everyone above a fixed score is hired).

2 Preliminaries

Notation. $N = \{1, \dots, n\}$ are the candidates, numbered in their (fixed) interview order. A symmetric matrix $W \in \{0,1\}^{n \times n}$ records friendships: $F(i)$ is the set of i's friends and $L(i) = \{j \in F(i) : j > i\}$ those who interview after i. Candidate i has a skill s_i and a coarse type $\theta_i \in \Theta = \{\text{weak}, \text{average}, \text{strong}\}$, the tercile of s_i. The company has K seats, a bank of Q questions, and asks m per interview. Y_j is j's score, H the hired set, and $x_j = \mathbf{1}[j \in H]$. The sharing matrix $S \in \{0,1\}^{n \times n}$ has $S_{ij} = 1$ if i tells j; row $S_{i\cdot}$ is i's action.

Definition 1 (Bayesian game and strategy)**.** A Bayesian game is a tuple $\langle N, (A_i), (\Theta_i), p, (u_i) \rangle$: players, action sets, type sets, a common prior p over types, and utilities u_i depending on all actions and types [11]. A strategy is a map $\sigma_i : \Theta_i \to \Delta(A_i)$; it is pure if it never randomises. *A strategy is a rule "given my type, tell these friends".*

Definition 2 (Bayes–Nash equilibrium, BNE)**.** A profile σ is a BNE if for every player i, type θ_i and action a_i', $\mathbb{E}[u_i(\sigma_i(\theta_i), \sigma_{-i}) \mid \theta_i] \geq \mathbb{E}[u_i(a_i', \sigma_{-i}) \mid \theta_i]$. The *exploitability* of σ is the largest expected gain any type can obtain by deviating; σ is an ε-BNE if its exploitability is at most ε. *Nobody, whatever their type, can do better by changing only their own rule.*

Definition 3 (Coarse correlated equilibrium, CCE)**.** A distribution z over strategy profiles is a CCE if for every player i and every fixed alternative σ_i', $\mathbb{E}_{\sigma\sim z}[U_i(\sigma)] \geq \mathbb{E}_{\sigma\sim z}[U_i(\sigma_i', \sigma_{-i})]$, where U_i is expected utility over types and the environment. Every BNE is a CCE, and the CCEs form a polytope, so linear programming finds them. *CCEs are what players who "have no regrets" can end up playing, so they are the right yardstick for learners.*

Definition 4 (Contextual bandit and regret)**.** In round t a learner sees a context x_t (here her type), picks an arm a_t (a row of S), and observes only the reward r_t of that arm. Its regret after T rounds is $\max_\pi \sum_{t\leq T} r_t(\pi(x_t)) - \sum_{t\leq T} r_t(a_t)$, the reward lost against the best fixed rule π from contexts to arms. *The learner never sees what the other choices would have paid.*

3 Problem Formulation

Stage game $\Gamma_{R,\phi,W,E}$, indexed by the hiring rule R, the utility type ϕ, the network W and the leak model E.

- *Order and types.* $s_i \overset{iid}{\sim} \mathcal{N}(0,1)$; θ_i is the tercile of s_i, and i observes only θ_i.
- *Action.* $S_{i\cdot} \in \{0,1\}^n$ with $S_{ij} = 0$ unless $j \in L(i)$: a candidate can tell only friends who interview after her. With two friend groups $\{1,3,5\}$ and $\{2,4,6\}$ exactly six entries are free: $S_{13}, S_{15}, S_{35}, S_{24}, S_{26}, S_{46}$.
- *Leaks.* Let $n_j = \sum_{i<j} S_{ij}$ be the number of people who tell j, and k_j the number of j's m questions seen in advance. **E0**: $k_j = m\,\mathbf{1}[n_j > 0]$ (a told question always comes up). **E1**: $k_j \sim \mathrm{Bin}\big(m, 1-(1-m/Q)^{n_j}\big)$ (comes up by chance). **E2**: $k_j = |A_j \cap \bigcup_{i:S_{ij}=1} A_i|$ with $A_j \sim \mathrm{Unif}\binom{[Q]}{m}$ (actual questions drawn from the bank).
- *Score.* $Y_j = s_j + \frac{1}{m}\big(\Delta k_j + \sum_{q=1}^{m} \eta_{jq}\big)$, $\eta_{jq} \sim \mathcal{N}(0,\sigma^2)$: skill, plus a bonus Δ per question seen in advance, plus marking noise. Candidates never observe Y_j.
- *Hiring.* SEATS: $H = \text{top-}K(Y)$. BAR: $H = \{j : Y_j \geq \tau_K\}$ with $\tau_K = \sqrt{1+\sigma^2/m}\;\Phi^{-1}(1-K/n)$, the cut-off that hires K people on average when nothing leaks.
- *Utilities*, with $|S_{i\cdot}|$ the number of friends told and c a small cost per tell: $u_i^{\text{SELF}} = x_i - c|S_{i\cdot}|$; $u_i^{\text{FRIEND}(\alpha)} = u_i^{\text{SELF}} + \alpha \sum_{j\in F(i)} x_j$; $u_i^{\text{CLAN}(\alpha,\beta)} = u_i^{\text{FRIEND}(\alpha)} - \beta \sum_{j\notin F(i)\cup\{i\}} x_j$; $u_i^{\text{FAIR}(a,b)} = u_i^{\text{SELF}} - \frac{1}{n-1}\sum_j \big[a(x_j-x_i)^+ + b(x_i-x_j)^+\big]$ [12]. *Selfish candidates want only their own offer; friendly ones also value a friend's offer at α of their own; clannish ones dislike outsiders' offers; fairness-minded ones dislike being behind (a) or ahead (b).*

Level 1 (known rules): find the equilibria of Γ. **Level 2 (unknown rules):** rounds $t = 1, \ldots, T$ are i.i.d. copies of Γ; candidate i knows her own utility function and her arms, but not E, (Q, m, Δ, σ) or anyone else's utility or strategy; she observes θ_i^t, chooses $S_{i\cdot}^t$ and sees only u_i^t.

Research questions. RQ1: the equilibrium $S^\star$ and the break-even friendliness $\alpha^\star_{i,\theta}$ (the smallest α at which type θ of candidate i tells), by position, type, W and R. RQ2: how utility types and mixed populations change it. RQ3: for each learner, regret, agreement with the equilibrium and CCE gap. RQ4: the company's hiring-value loss $\Lambda_R = \mathbb{E}[V \mid \text{no leaks}] - \mathbb{E}[V]$, $V = \sum_{j\in H}(s_j - \tau_K)$, as a function of Q.

Structural results. Let ℓ_i be i's own loss in hiring probability from sharing, and G_i the gain in her friends' expected hires.

Lemma 1. *With all random draws and all other rows fixed, setting $S_{ij} = 1$ weakly raises Y_j and changes no other score. Hence $\ell_i \geq 0$ and $G_i \geq 0$ under both rules.*

Proposition 1 (No sharing in a clique)**.** *With SEATS, complete W, FRIEND(α), $\alpha \leq 1$ and $c > 0$, no one shares in any BNE.*

Proof. $|H| = K$, so $|F(i) \cap H| = K - x_i$ and $u_i = \alpha K + (1-\alpha)x_i - c|S_{i\cdot}|$. Sharing weakly lowers x_i (Lemma 1) and costs c. □

Corollary 1 (The popular candidate stays silent)**.** *The proof uses only $F(i) = N \setminus \{i\}$. So under SEATS a candidate who is friends with* everyone *never shares, whatever the rest of the network looks like and whatever the others do.*

Proposition 2 (Sharing is free under a cut-off)**.** *Under BAR and $\phi \in \{\text{SELF}, \text{FRIEND}, \text{CLAN}\}$, $\ell_i \equiv 0$ and the gain from any row does not depend on θ_i, so each candidate has a best response that is the same for all her types. (For FAIR this can fail: the penalty depends on whether i herself is hired.)*

Proposition 3 (Leak formula)**.** *In E2, each of j's questions is leaked with probability exactly $1-(1-m/Q)^{n_j}$. The expected leaked fraction is at most ε if $Q \geq m/(1-(1-\varepsilon)^{1/n_j})$, and $Q \geq mn_j/\varepsilon$ always suffices.*

Proofs. Lemma 1: Y_j is monotone in k_j, and k_j in n_j. Prop. 2: $x_i = \mathbf{1}[Y_i \geq \tau_K]$ does not depend on S. Prop. 3: independent uniform sets and $(1-\varepsilon)^{1/n} \leq 1 - \varepsilon/n$. *Remark (clans).* Under SEATS, $|H| = K$ gives $u_i^{\text{CLAN}(\alpha,\beta)} = (1+\beta)x_i + (\alpha+\beta)\sum_{F(i)} x_j - \beta K - c|S_i.|$: up to a constant, a friendly candidate with $\alpha_{\text{eff}} = (\alpha+\beta)/(1+\beta) > \alpha$. Disliking outsiders should therefore *raise* sharing under SEATS and do nothing under BAR.

4 Methodology

The pipeline has three stages: simulate the game into a table, solve the table (Level 1), and let learners play the simulator (Level 2). All experiments use $n = 6$, $K = 2$, $m = 3$, $Q = 15$, $\Delta = \sigma = 1$, $c = 0.001$ unless a parameter is being swept.

4.1 Stage 1: from simulator to payoff table

For each of the $3^6 = 729$ type profiles θ and each of the 2^f sharing matrices S (f free entries; $f = 6$, i.e. 64 matrices, for two friend groups) we simulate 2,000 interview days and store $P(x_j = 1 \mid \theta, S)$, $P(x_i x_j = 1 \mid \theta, S)$ and $\mathbb{E}[V \mid \theta, S]$ for both hiring rules. Three choices make this cheap and precise:

- *Same days for every S.* Skills, marking noise and leak draws are generated once per day and reused for all matrices and both rules (common random numbers). Payoff *differences* between two sharing choices — all that equilibria depend on — are then far more precise than the payoffs themselves, and Lemma 1 holds exactly in the table, not just on average.
- *Stratify by type.* Skills are drawn from the normal distribution truncated to each candidate's tercile, so all 729 profiles get the same sample size.
- *Simulate once, re-use for every utility.* All four utility families are linear in the stored probabilities (FAIR needs the pairwise ones, since $(x_j - x_i)^+ = x_j - x_i x_j$). Changing α, β, (a, b) or mixing utility types needs no new simulation.

One table for two friend groups (93 million simulated outcomes) takes 15–40 s on a laptop. We also store the means over 20 blocks of 100 days, used for the bootstrap below.

4.2 Stage 2 (Level 1): solving the known game

A pure strategy of candidate i assigns one of her $2^{|L(i)|}$ rows to each of her three types. With two friend groups this gives $4^3 \cdot 4^3 \cdot 2^3 \cdot 2^3 = 262{,}144$ profiles. The central object is the *interim value*

$$V_i[\sigma_{-i}, \theta_i, a] = \mathbb{E}_{\theta_{-i}}\big[u_i\big(\theta,\, S(a, \sigma_{-i}(\theta_{-i}))\big)\big],$$

the expected utility of type θ_i of candidate i from row a when the others follow σ_{-i}. We precompute it for every σ_{-i}; everything else is a lookup.

- **Pure equilibria: exhaustive.** For every profile we compute the exploitability $\max_{i,\theta_i}\big(\max_a V_i[\sigma_{-i}, \theta_i, a] - V_i[\sigma_{-i}, \theta_i, \sigma_i(\theta_i)]\big)$ and keep the profiles with value 0. The list of pure BNEs is therefore complete, not a sample from a search.
- **Independent checks.** Best-response dynamics from 8 starting profiles (nobody shares, everybody shares, 6 random) must end in an enumerated equilibrium. To look for mixed equilibria we run replicator dynamics on the 12 (candidate, type) agents from 3 random interior starts.
- **Mixed equilibria: dominance and support enumeration.** Replicator dynamics only finds *stable* mixed equilibria, so on its own it says little about whether others exist. We therefore apply iterated strict dominance to every game: an action that loses against every surviving pure profile of the others also loses against every mixed profile, so no equilibrium (pure or mixed) uses it. If every (candidate, type) keeps a single action, that profile is the game's only equilibrium of any kind, and it is checked to be an exact pure equilibrium of the table. On the pairs network, where every agent has two actions, we then enumerate the supports of what dominance leaves and solve the indifference conditions (exactly when at most two agents mix; by multi-start root finding, which is *not* exhaustive, when three do), checking every candidate against all deviations of all agents. On the other two networks we run only the dominance step.
- **All CCEs at once.** A CCE is a distribution over the profiles satisfying one linear inequality per (candidate, alternative strategy) pair. We solve two linear programs (HiGHS) for each of the three leak models: maximise and minimise the expected share of tells. If the two optima coincide, *every* CCE — hence every mixed equilibrium and every limit of no-regret learning — produces the same amount of sharing.
- **How sure are we? Bootstrap ε.** The table is an estimate. We resample its 20 blocks of days with replacement (200 times; 100 for the two smaller networks), rebuild the utilities (resampling the table of pairwise hiring probabilities with the same blocks when a utility needs it, as the fairness utility does), and recompute the exploitability of the reported equilibrium. We report the share of replicates in which it is still an exact equilibrium, and the 95th percentile of ε.

Algorithm 1: One round of Leak-aware UCB for candidate i of type θ (arms $a \subseteq L(i)$)

play $a = \arg\max_{a'} \ \hat{\mu}_{\theta,a'} + \sqrt{2 \ln t_\theta / n_{\theta,a'}}$; observe own utility r and, for each told friend $j \in a$, the number $k_j^{(i)}$ of told questions that came up;

$D \leftarrow \{j \in a : k_j^{(i)} = 0\}$; `// tells that did not land`

for *every* a' *with* $a \setminus D \subseteq a' \subseteq a$ **do**

  $n_{\theta,a'} \leftarrow n_{\theta,a'} + 1$; $\hat{\mu}_{\theta,a'} \leftarrow$ running mean with the new sample $r + c\,|a \setminus a'|$; `// refund cost of unmade tells`

log leak-rate estimate $\hat{p} \leftarrow$ mean of $k_j^{(i)}/m$ over all tells so far (reporting only, not used to choose arms); $t_\theta \leftarrow t_\theta + 1$;

- **Break-even friendliness** is measured twice: by the formula $\alpha^\star = (\ell_i + c)/G_i$ from the table, and directly as the smallest α on a grid of step 0.025 at which that type tells in every equilibrium.

The complete network has 15 free entries (2^{15} matrices; too many profiles to enumerate at $n = 6$). There we solve $n = 4$ exhaustively and, at $n = 6$, test the two candidate profiles "nobody shares" and "everybody shares" against every deviation of every type of every candidate.

4.3 Stage 3 (Level 2): learning the unknown game

Each (candidate, type) pair is one bandit learner whose arms are that candidate's rows; with two friend groups there are 12 learners with 4 or 2 arms. Each round draws fresh skills, the realised types' learners choose, the simulator runs one day, and each learner receives only its own realised utility (rescaled to $[0, 1]$). We compare six learners:

- **ε-greedy**: explore uniformly with probability $\varepsilon_t = t^{-1/3}$, else play the best arm so far. **UCB1** [6]: play $\arg\max_a \hat{\mu}_a + \sqrt{2 \ln t / n_a}$. **Thompson sampling** [8]: sample each arm's mean from a Gaussian posterior, play the largest. **EXP3** [7]: exponential weights on importance-weighted rewards; assumes nothing about how rewards are generated.
- **Leak-aware UCB** (ours; Algorithm 1). The review promised a "model-based" learner that estimates the leak rate and Δ and best-responds in its estimated game. Best-responding needs the other candidates' strategies, which nobody observes, so we split the idea in two. This first learner is the cheap half: UCB1 plus one fact about how leaks work, with no model of the environment that it plans with (the leak-rate estimate in Algorithm 1 is logged but never used to choose an arm). The fact is: *a tell matters only through questions that actually come up.* If my friend reports that none of the questions I told her appeared, the day went exactly as if I had not told her, so the same observation is also a valid sample for the arm "did not tell her" (plus the small cost c I would have saved). In E1/E2 about half of all tells do not land ($0.8^3 = 0.51$), so most pulls of a "tell" arm also teach about the "silent" arms. The price is one extra piece of feedback: how many of my questions came up for my friend. Under E0 every tell lands, and the learner is identical to UCB1. In E1 the credited sample has exactly the distribution of a sample of the shorter arm; in E2 the day is identical, but keeping only days on which the two question sets are disjoint slightly changes the chance that other tellers' questions land, so there the credit can be biased. We measured it in the worst case (candidate 5 has a second teller; 4 million days): the bias is at most 0.08 percentage points, about two standard errors (0.03), which is 22% of the smallest payoff gap the learners must detect (0.36 points), and a control with no second teller scatters as much (up to 0.06); so we cannot show it is zero, only that it is small next to the signal, and agreement is 100% in E2.
- **Model-based UCB** (ours). The other half: a learner that estimates how leaks work and plans with the estimate. For each (candidate, type) it fits, by ridge regression on *every* round whatever arm was played, the additive model $r + c_n|a| = \mu_0 + \sum_{j \in a} \beta_j k_j + \text{noise}$, where k_j is the number of the questions it told friend j that came up (the same feedback as above, zero for friends not told), c_n is the known cost of a tell in reward units, μ_0 is the value of keeping quiet and β_j the value to the teller of one leaked question reaching j. It also estimates the leak rate $\hat{p}$ from all observed k_j/m, and plays the arm with the highest predicted value $\mu_0 + \sum_{j \in a} \beta_j m\hat{p} - c_n|a|$ plus a LinUCB bonus. The bonus is computed with an optimistic leak rate ($\hat{p}$ plus a Hoeffding radius): with the plug-in $\hat{p}$ alone, one tell that happens not to land gives $\hat{p} = 0$, hence no exploration bonus on β, and the arm is never tried again (a unit test caught this). This is a reduced-form version of the learner in the review: β_j is Δ passed through the hiring rule and the others' play rather than Δ itself, planning is single-agent rather than a best response in an estimated game of all candidates, and the model is additive across friends, so it cannot represent that a second teller of the same friend is redundant (E0).

What we measure. *Agreement*: the share of the 12 learners whose currently-best arm is the equilibrium row (if a game has several equilibria, the best-matching one). *Regret*: Definition 4, computed with the table's expected payoffs so that luck does not enter. *CCE gap*: the largest average gain any candidate could have made by switching to one fixed rule for the whole run (distance of the history of play from a CCE). *Rounds to 90%*: first round at which agreement reaches 0.9. *Residual gain*: the exploitability of the profile of most-played arms — what anyone could still gain by deviating from where the learners settled.

Table 1: Verification checks. "Deviations" are all rows of every type of every candidate, against both "nobody shares" and "everybody shares".

Check	What must hold	Simulator	Verdict
Worked example	0.7083, 0.9149, 0.7083 (exact)	0.7085, 0.9157, 0.7088	pass
Lemma 1	a tell never lowers the friend's score, changes no one else's	min change 0, others 0	pass
Prop. 3	leaked fraction $= 1-(1-m/Q)^{n_j}$, $n_j = 1,2,3$	largest error 0.0013 (2.2 s.e.)	pass
Prop. 2	under a cut-off, own row never changes own hire	largest change 0	pass
	... and friends' gain is the same for every own type	largest spread 0.0029	pass
Prop. 1	everyone friends, fixed seats: nobody tells (n=4 solved, n=6 deviations)	unique equilibrium; 0 of 684 deviations pay	pass
Selfish baseline	nobody tells at $\alpha = 0$ (both rules, E0–E2)	6 of 6 games	pass

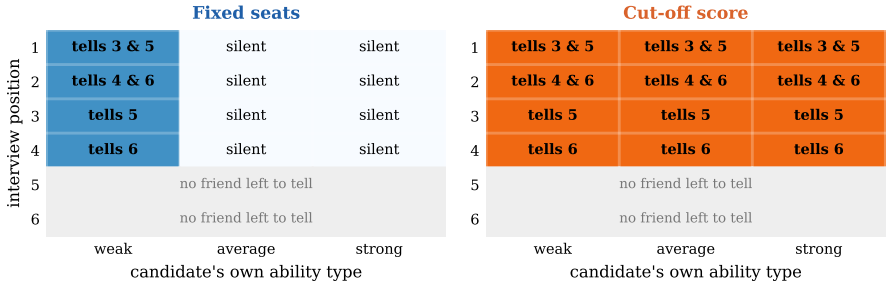


Figure 1: The equilibrium for two friend groups $\{1,3,5\}$, $\{2,4,6\}$ at friendliness $\alpha = 0.2$ (leak model E1). Under either rule it is the only pure equilibrium among 262,144 profiles.

Experiments and statistics. X1: rule $\times$ $\alpha \in \{0, 0.025, \ldots, 0.5\}$ $\times$ four networks $\times$ E0–E2. X2: four utility families and five mixed populations. X3: six learners $\times$ two rules $\times$ E0–E2, $T = 200{,}000$ rounds, 30 seeds each. X4: bank size $Q \in \{4, \ldots, 64\}$ in E2, equilibrium re-solved for each Q. Runs that share a seed share every random draw of the environment across rules, leak models and learners, so comparisons are paired. We report 95% bootstrap confidence intervals over seeds, two-sided sign-flip permutation tests on the paired differences, and Holm-corrected p-values within each family of tests [13].

Checks before any experiment (Table 1). The simulator must reproduce the worked example and every structural result of Section 3; all pass. A further 42 unit tests cover every module: the simulator and tables, the game and the solvers (including replicator dynamics, the CCE program and the bootstrap), support enumeration (against matching pennies, a coordination game and a dominance-solvable game), the learners, the learning loop (including that the leak-aware learner is identical to UCB1 under E0), the statistics, the parameter sweep, the pivotality and baseline calculations.

5 Results

5.1 Who tells? The hiring rule decides (RQ1)

Under fixed seats only weak candidates tell; under a cut-off everyone does (Figure 1). The pilot's picture survives exhaustive search: with fixed seats 33% of the possible tells are made, with a cut-off 100%. Under both rules the equilibrium is unique, all 8 best-response runs and all 3 replicator runs end in it, and it remains an exact equilibrium in 100% (seats) and 100% (cut-off) of 200 bootstrap replicates of the table.

Why: a weak candidate has almost nothing to lose (Figure 2a,b). Under fixed seats, telling one friend costs a weak candidate 0.38 percentage points of her own hiring chance, an average one 1.49 and a strong one 1.60 — a weak candidate is rarely near the last seat, so a boosted friend rarely pushes *her* out. This does not rest on the loss itself: measured on fresh days with nobody sharing, a weak candidate holds the last seat on 3.6% of days, an average one on 17.9% and a strong one on 28.5%. The friend gains about the same in all three cases (4.5, 4.9, 4.5 points). The break-even friendliness $\alpha^\star = (\ell + c)/G$ is therefore 0.10, 0.32 and 0.38 for weak, average and strong candidates (the pilot estimated 0.09, 0.32, 0.37), and the exact equilibria switch at the next grid points: 0.125, 0.325–0.35 and 0.375–0.4. Under a cut-off the loss is exactly zero for every type (Proposition 2) and $\alpha^\star \approx 0.01$: any friendliness above the cost of a message is enough.

The staircase holds across networks and leak models (Figure 3). We solved 378 games exactly. Under a cut-off, sharing jumps from nothing to its maximum at the first grid point above zero in every network. Under fixed seats it climbs in steps — weak, then average, then strong. A cut-off produces at least as much sharing as fixed seats in 188 of 189 comparisons, and strictly more in 134; the one exception is discussed below. Three details:

- *The popular candidate never tells.* In the star network the hub is friends with everyone. Her break-even friendliness is above 1 and she is silent at every $\alpha \le 0.5$ under fixed seats, as Corollary 1 says: any seat she hands a friend is taken from herself or from another friend. Sharing in the star therefore stops at 40% (the two candidates before her).

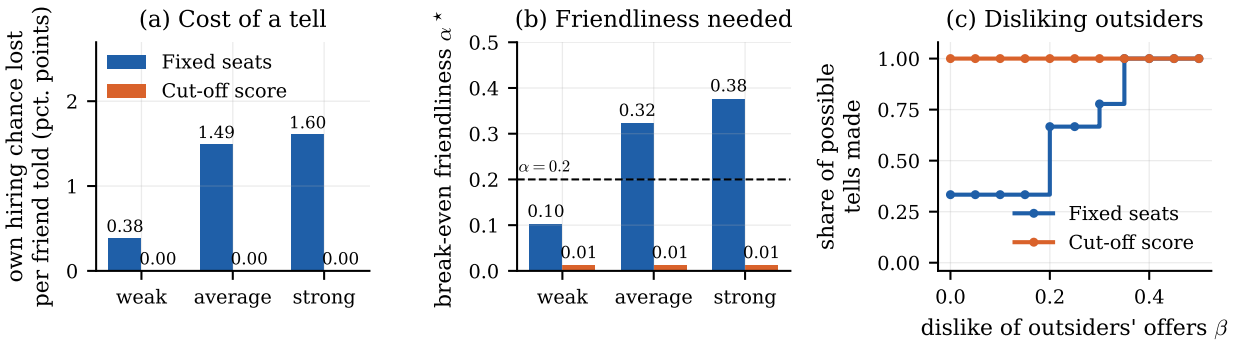


Figure 2: (a) Own hiring chance lost per friend told. (b) Break-even friendliness; types below the dashed line tell at $\alpha = 0.2$. (c) Share of possible tells made in equilibrium as candidates dislike outsiders' offers more ($\alpha = 0.2$). Two friend groups, E1.

- *Everyone friends: nobody tells* under fixed seats ("nobody shares" passes the deviation test at all 21 values of α, "everybody shares" at 0), and everybody tells under a cut-off (at 19 of the 20 values $\alpha > 0$; at the smallest, a few tells are not worth the message cost). This is Proposition 1 at $n = 6$.
- *When told questions always come up (E0), a second tell is wasted.* A candidate needs to be told only once, so equilibria differ in *who* tells her, and sharing tops out below 100% (4 of 6 tells in two groups). This is the only place we find several equilibria: 63 of 378 games, all in E0; the other 315 have exactly one and none has zero. It is also where the one exception sits (two friend groups, E0, $\alpha = 0.375$): fixed seats have 8 equilibria sharing 61–78%, the cut-off 4 sharing 67%.

How solid are these equilibria? The reported profile stays an exact equilibrium in at least 95% of bootstrap replicates in 370 of 378 games (98%). The exceptions sit on a step of the staircase, where a type is nearly indifferent; even there the 95th-percentile ε is at most 0.03 percentage points of utility. So the *pattern* is certain and each *threshold* is known to within one grid step. Uniqueness is a statement about the estimated table, so we also count near-ties: profiles that are within $\varepsilon = 0.001$ (one message cost, 0.1 percentage points) of being an equilibrium. The games of Figure 1 ($\alpha = 0.2$) still have exactly 1 (seats) and 1 (cut-off). Over all 378 games, 117 have more than one: 17 at $\alpha \leq 0.025$, where a tell is worth about what it costs; 62 in E0, where a redundant tell costs exactly c; and 38 on the steps of the fixed-seats staircase in E1/E2 (up to 64 profiles), where one type is nearly indifferent. Away from these, every cut-off game in E1/E2 has a single near-equilibrium. Best-response dynamics converged in 3024 of 3024 runs, and replicator dynamics ended at an enumerated pure equilibrium in 375 of 378 runs (the other 3, all in the most fragile game, had not settled after 3,000 steps).

What about mixed and correlated equilibria? The CCE linear programs bound what any mixed or correlated equilibrium can do. In E1 and E2 the maximum and minimum sharing over *all* CCEs coincide in 106 of 106 (E1) and 106 of 106 (E2) games, so there no mixed or correlated equilibrium shares more or less than the pure one. E0 is the exception: the interval is a single point in only 55 of 106 games, and in 48 of them it is wider than the range spanned by the pure equilibria (up to 42 percentage points wide). For two friend groups at $\alpha = 0.2$ under a cut-off, the 4 pure equilibria all share 67% of the possible tells, while correlated play can share anywhere from 66% to 99%. A plausible reason, which we did not test: in E0 a second tell of the same friend is exactly redundant and costs only c, so candidates are nearly indifferent about who tells and a CCE can spread redundant tells almost freely. The pure-equilibrium statements of this section therefore describe E0 less completely than E1 and E2. Iterated strict dominance (Section 4) settles the rest: it solves 311 of the 378 games (pairs 123/126, star 96/126, two friend groups 92/126), and in each of them the surviving profile is the game's only equilibrium of any kind. All but 1 of the games it leaves open are in E0 (the exception: two friend groups, E1, fixed seats, $\alpha = 0.1$). On the pairs network, support enumeration finds exactly one mixed equilibrium in each of the 3 games dominance leaves open (E0, fixed seats), next to the two pure ones, with sharing between theirs and inside the CCE interval; since three agents mix there and the root finding is not exhaustive, there may be more. On the other two networks nothing is claimed about mixed equilibria beyond the CCE bound.

5.2 What candidates care about (RQ2)

Disliking outsiders makes candidates more generous to friends — but only under fixed seats (Figure 2c). As β rises from 0 to 0.5 at fixed $\alpha = 0.2$, sharing under fixed seats rises from 33% to 100% (everyone tells from $\beta = 0.35$, i.e. $\alpha_{\text{eff}} = 0.41$, in line with the strong types' threshold above), and stays at 100% under a cut-off. This is the Remark of Section 3: with a fixed number of seats, each seat a friend takes is one an outsider loses.

Fairness-minded candidates behave the other way round (Table 2). Under fixed seats they never tell. Under a cut-off it is the *strong* who tell, and only for moderate inequality aversion: a strong candidate expects to be hired and dislikes being ahead of her friends, so she pulls them up; if envy (a) is large, the risk of being the one left behind

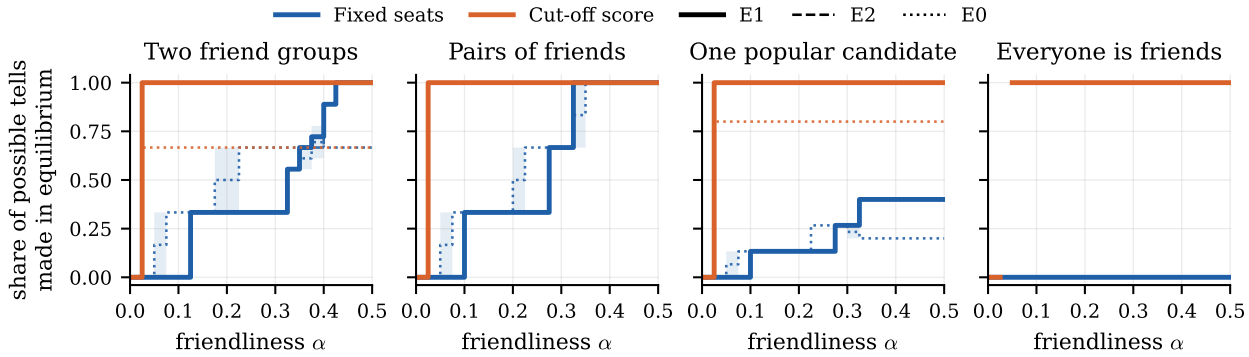


Figure 3: Share of possible tells made in equilibrium as candidates care more about friends. Shaded bands: range over multiple equilibria (E0 only). "Everyone is friends" shows the profile that passes the deviation test (E1).

Table 2: Equilibrium sharing for different motives (two friend groups A= $\{1,3,5\}$, B= $\{2,4,6\}$; E1). Every game has exactly one pure equilibrium. †: exact equilibrium in fewer than 95% of bootstrap replicates (a type is nearly indifferent).

What candidates care about	Who tells, in equilibrium		Tells made (%)	
	Fixed seats	Cut-off	Seats	Cut-off
Only own offer (selfish)	nobody	nobody	0	0
Friends' offers too, α=0.2	weak	everybody	33	100
... and dislike outsiders, β=0.2	weak; average†	everybody	67	100
... and dislike outsiders, β=0.4	everybody	everybody	100	100
Dislike inequality, (a,b)=$(0.25, 0.1)$	nobody	nobody	0	0
Dislike inequality, (a,b)=$(0.5, 0.25)$	nobody	strong (4 of 6 tells)†	0	22
Dislike inequality, (a,b)=$(1, 0.5)$	nobody	strong	0	33
Dislike inequality, (a,b)=$(2, 0.5)$	nobody	nobody	0	0
Group A friendly, group B selfish	weak in A	everyone in A	17	50
Only the first interviewees friendly	weak (cand. 1–2)	everyone (cand. 1–2)	22	67
Only the first interviewees selfish	weak (cand. 3–4)	everyone (cand. 3–4)	11	33
Group A friendly, group B clannish (β=0.3)	weak; average in B; strong (2 of 6 tells)†	everybody	61	100
Group A inequality-averse, group B friendly	weak in B	weak in B; average in B; strong (5 of 6 tells)	17	61

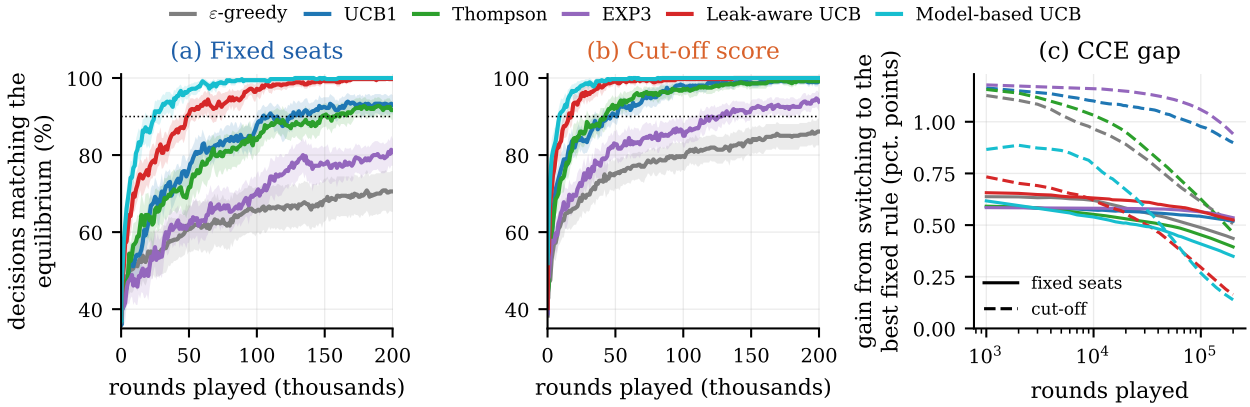


Figure 4: Learning in E1, mean over 30 seeds with 95% bootstrap bands. (a,b) Share of the 12 (candidate, type) learners whose best arm is the equilibrium row. (c) CCE gap of the history of play.

dominates and telling stops. This is the case Proposition 2 excludes, and the only one where behaviour under a cut-off depends on type.

Motives barely spill over. In the 10 mixed-population games, only 1 of 120 (candidate, type) choices differs from what that candidate does in a population made entirely of her own kind (a nearly indifferent strong clannish candidate tells one friend more when the other group is merely friendly). With one friendly and one selfish group, exactly the friendly group shares: weak types under fixed seats, everyone under a cut-off.

5.3 Can candidates learn it? (RQ3)

Yes — but the signal is tiny, so it takes tens of thousands of interviews. At equilibrium the best row beats the second-best by only 0.36–0.80 percentage points of utility under fixed seats and 0.96–0.99 under a cut-off, while a single outcome is 0 or 1. UCB1 and Thompson sampling nevertheless end with 93% and 92% (seats) and 99% and 99% (cut-off) of decisions correct in E1 (Table 3); EXP3, which guards against adversarial rewards, and ε-greedy are slower (81/94% and 71/86%). Results under E2 are the same within error bars.

Table 3: Decisions matching the equilibrium after 200,000 rounds (%, mean ± half-width of the 95% CI over 30 seeds; best per column in bold), and median rounds to reach 90% ("never (x%)": the median run did not, x% of runs did; E0 and E2 in Table 5), and cumulative regret after 200,000 rounds (Definition 4, in units of one offer, mean over the 12 learners and 30 seeds; lowest in bold). In E0 with a cut-off the equilibria differ only in who makes a redundant tell.

	E0		E1		E2		Rounds to 90% (E1)		Regret at T (E1)	
Learner	Seats	Cut-off	Seats	Cut-off	Seats	Cut-off	Seats	Cut-off	Seats	Cut-off
ε-greedy	82 ±4	**74 ±3**	71 ±5	86 ±3	72 ±4	86 ±4	never (20%)	never (50%)	**154**	178
UCB1	**96 ±2**	56 ±2	93 ±3	99 ±1	94 ±2	99 ±1	56k	18k	246	421
Thompson sampling	93 ±2	70 ±4	92 ±3	99 ±1	94 ±2	99 ±1	78k	19k	173	191
EXP3	93 ±2	51 ±1	81 ±4	94 ±2	84 ±4	93 ±2	141k	61k	251	442
Leak-aware UCB	**96 ±2**	56 ±2	100 ±0	**100 ±0**	**100 ±0**	**100 ±0**	27k	7k	247	70
Model-based UCB	**96 ±2**	53 ±2	**100 ±0**	**100 ±0**	**100 ±0**	**100 ±0**	12k	4.5k	168	**65**

Finding the equilibrium and earning along the way are different things. Regret ranks the learners almost the other way round (Table 3, last two columns). Under fixed seats in E1, cumulative regret is 154 for ε-greedy and 173 for Thompson sampling, against 246 for UCB1 and 251 for EXP3; under a cut-off 178 and 191 against 421 and 442. So UCB1 ends up with the right answer but pays more to get there, and ε-greedy has the lowest regret under fixed seats despite the lowest agreement. We checked where the regret comes from. UCB1 puts 56% of its pulls on non-equilibrium rows under fixed seats and 48% under a cut-off (Thompson sampling: 41% and 23%; ε-greedy: 37% and 23%): with payoff gaps this small its confidence bonus never lets it settle, so it ends with the right most-played row but keeps paying for the others. ε-greedy's low regret is *not* because its mistakes are on cheap decisions (the mean payoff gap of the cells where it ends wrong is 1.01 times the mean over all cells; we had guessed otherwise) but because it pulls wrong rows less often, while agreement only looks at the row that ends up on top. The regret implied by these pull counts matches the stored regret to within 6% (the leak-aware learner is left out: its counts include credited observations).

A cut-off is easier to learn than fixed seats whenever the equilibrium is unique: the gap is larger and the same for all types. In E1 and E2, agreement under a cut-off is at least as high for every learner and significantly higher (Holm-corrected $p < 0.05$) in 8 of 12 paired tests (the exceptions: Leak-aware UCB, Model-based UCB, at 100% under both rules). Under fixed seats the slowest learners are the *average* types (UCB1: 99% of weak, 83% of average, 96% of strong decisions correct), although their gap is not the smallest (0.40, 0.55, 0.76 points). The reason is noise: at the equilibrium the reward of a weak type has standard deviation 0.15 (rewards lie in $[0, 1]$), against 0.30 for average and 0.31 for strong types, because a weak candidate is almost never hired. Dividing squared noise by the squared payoff gap, which is what a bandit has to average away, gives 1,474, 2,934 and 1,730 for weak, average and strong types, so the average types face the hardest problem; over the 12 decision cells this ratio is rank-correlated with the final error rate ($\rho = 0.80$ for UCB1 and 0.70 for Thompson sampling).

Knowing how leaks work is worth a factor of two to five. Two learners use it. Leak-aware UCB (UCB1 plus credit for tells that did not land) reaches 90% agreement in E1 after 27,000 rounds under fixed seats and 7,000 under a cut-off, against 56,000 and 18,000 for UCB1. Model-based UCB, which estimates the leak rate and the value of a leaked question and plans with them, needs 12,000 and 4,500. Both end at 100% (rounded) in E1 and E2 under both rules. At 20,000 rounds Leak-aware UCB is ahead of UCB1 by +13 and +12 points (E1) and +11 and +7 points (E2), and Model-based UCB by +24 and +15 (E1) and +20 and +10 (E2); all 4 and all 4 differences are significant after Holm correction. The settled profile of Model-based UCB is an exact equilibrium in 100% (seats) and 100% (cut-off) of E1 runs, of Leak-aware UCB in 97% and 100%, and of UCB1 in 40% and 87%. The two help in different ways. Against UCB1, the cumulative regret of Leak-aware UCB changes by −351 under a cut-off, where the best arm is "tell everyone" and each pull also informs every other arm, and by +2 under fixed seats, where the best arm for most types is silence — and silence teaches nothing about telling. Model-based UCB lowers regret in all four E1/E2 settings (−77 and −356 in E1, −67 and −330 in E2; significant after Holm correction in 4 of 4), and it is the only learner that improves on UCB1 in both regret and agreement under fixed seats (ε-greedy and Thompson sampling also have lower regret there, but lower agreement). This is consistent with pooling information across rows that tell the same friend: its relative regret reduction against UCB1 is larger in the four-row cells (candidates with two later friends) than in the two-row cells in 4 of 4 settings (E1, fixed seats: 34% against 25%; cut-off: 86% against 81%), though the difference is small under a cut-off, and the four-row cells differ in other ways too. Its leak-rate estimate is used to choose arms, and is right ($\hat{p} = 0.200$ in E1 and E2, truth $m/Q = 0.2$; 1.000 in E0, where every told question comes up). The model is additive across friends, so it cannot represent a redundant second tell. In E0 under a cut-off, where that matters, its agreement is 53% (UCB1: 56%) — see the caution below — while its regret is still lower (−96).

Play moves toward the CCE set, but is not yet close to it (Figure 4c). The CCE gap falls between 10,000 and 200,000 rounds for all six learners (to 0.42–0.73 of its earlier value; significant after Holm correction in 36 of 36 learner-by-setting tests, each with the 30 seeds as the unit; 24 of 24 in E1 and E2) and ends at 0.38–0.94 percentage

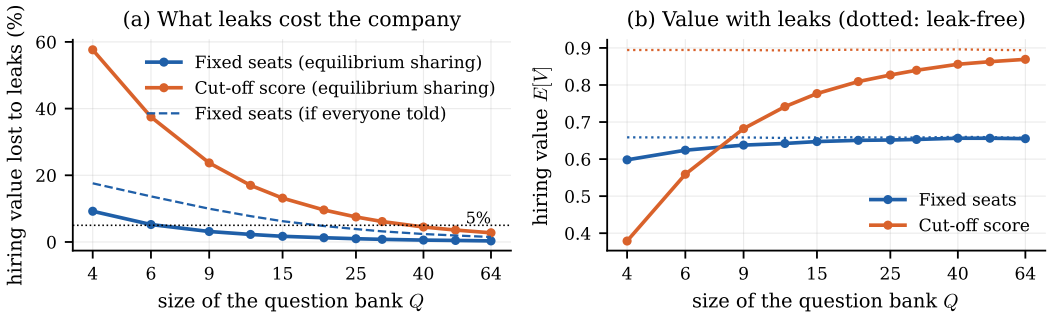


Figure 5: Company's hiring value under equilibrium sharing as the question bank grows (E2, $\alpha = 0.2$, $m = 3$ questions per interview). 95% block-bootstrap intervals (500 replicates) are narrower than the markers.

points. That is about the size of the payoff gaps the learners have to detect (0.36–0.80 points under fixed seats), so after 200,000 rounds the history of play is measurably closer to a CCE but not yet near one: exploration is still being paid for.

What the settled play is worth. Against the baselines of the review (Table 6, E1, $\alpha = 0.2$; E2 is the same within a few thousandths), no learner's settled profile is below random sharing in welfare, in any setting (0 exceptions in 24 comparisons). Under a cut-off, where the number of hires moves with sharing, welfare is 3.067 for random sharing and 2.799 when nobody shares, and the learners settle between 3.256 and 3.317, the equilibrium value. Under fixed seats exactly K candidates are hired whatever anyone shares, so welfare is the same for every pattern up to the message cost and the comparison is about the company: the equilibrium is worth 0.645 to the company, above random sharing (0.631) and everybody sharing (0.614), below nobody sharing (0.657, the leak-free company), and the learners settle between 0.638 and 0.645. Under a cut-off the picture for the company is the reverse and is the cost of leaks of Section 5.4, not a failure of learning: the equilibrium itself is worth 0.771 to the company, less than random sharing (0.833), so every learner is below random sharing there.

A caution about "agreement". In E0 with a cut-off, agreement looks poor (51–74%). The game there has 4 equilibria that differ only in which of two friends tells the last candidate; the learners end up both telling. The residual gain from deviating is 0.10 percentage points for every learner — exactly the message cost c. They have learned everything that matters and failed only to coordinate on who saves a message.

5.4 What it costs the company (RQ4)

At $\alpha = 0.2$, leaks hurt a cut-off company about ten times more — mostly because more people tell. With the pilot's bank of 15 questions, equilibrium sharing costs 1.7% of hiring value under fixed seats and 13.1% under a cut-off; across all bank sizes the cut-off loses 8–12 times more, and the bootstrap intervals are disjoint at all 11 sizes (Figure 5a). To keep the loss below 5% the company needs $Q \geq 9$ under fixed seats and $Q \geq 40$ under a cut-off. Proposition 3 explains the shape: the leaked fraction falls like $m\, n_j/Q$. Two effects add up, and they are not equally large. The first is *who* tells: at $\alpha = 0.2$, 100% of possible tells are made under a cut-off and 33% under fixed seats. The second is the rule itself: every leaked mark can lift someone over a fixed line, whereas under fixed seats a leak only reorders candidates. The dashed line in Figure 5a separates them. If everyone told under both rules, fixed seats would lose 6.2% at $Q = 15$ (against 13.1%), the cut-off would lose only 2.5–4.4 times more, and fixed seats would need $Q \geq 20$ rather than 9. Random sharing (every possible tell with probability $\frac{1}{2}$, the baseline of the review) would cost 3.6% under fixed seats and 6.5% under a cut-off at $Q = 15$: the equilibrium is better for a fixed-seats company than random sharing, and about 2.0 times worse for a cut-off company. The tenfold gap is therefore specific to a friendliness at which only weak candidates tell under fixed seats; wherever candidates are friendly enough that everyone tells under both rules, the everyone-tells comparison is the relevant one.

But the cut-off is still the better rule once the bank is large enough (Figure 5b). Without leaks a cut-off hires better (value 0.89 against 0.66): it can hire three good candidates or none, where fixed seats must always hire two. Leaks eat this advantage only for small banks: fixed seats give more value up to $Q = 6$, the cut-off from $Q = 9$ (0.78 against 0.65 at $Q = 15$). For a company the message is not "avoid cut-offs" but "a cut-off needs a question bank several times larger".

5.5 Verdict on the hypotheses, deviations and limits

Judged by the refutation criteria fixed in the review (Table 4), H1, H3 and H5 are supported. H4 is not refuted, with unequal support: the faster-learner half holds clearly (for a reduced-form model-based learner and for the simpler leak-aware one), but the CCE gap falls without becoming small. H2 is refuted at one point: in E0, where the game has several equilibria, some fixed-seats equilibria share more than the cut-off's; in E1 and E2 it holds in every comparison. **Deviations from the review plan:** (i) tables use 2,000 days per cell, not the pilot's 8,000; precision

Table 4: The five hypotheses fixed in the project review, the refutation criterion fixed with them, and what the experiments say.

	Hypothesis	Refuted if (review)	Evidence	Verdict
H1	Under fixed seats, $\alpha^\star$ is highest for the types most likely to be pivotal	$\alpha^\star$ flat, or lowest for pivotal types	$\alpha^\star = 0.10 < 0.32 < 0.38$, in the order of how often a type holds the last seat when nobody shares (3.6, 17.9, 28.5% of days, fresh simulation without sharing) and of the own loss ($\ell = 0.38, 1.49, 1.60$ points)	Supported
H2	Cut-off shares more than fixed seats for every $\alpha \in [0, 0.5]$	Fixed seats share more at some α	Holds in all 126 E1/E2 comparisons. In E0 fails once (two friend groups, E0, $\alpha = 0.375$): fixed seats have equilibria sharing up to 78%, the cut-off 67%	Refuted once (E0); holds in E1/E2
H3	Fixed seats: sharing rises with β; cut-off: no effect	No rise under fixed seats, or an effect under a cut-off	33%→100% (never falling) vs. constant 100%	Supported
H4	No-regret play nears the CCE set; model-based learners are faster	CCE gap does not fall, or not faster	CCE gap falls significantly in 36 of 36 learner-by-setting tests but ends at 0.38–0.94 points, the size of the payoff gaps; Model-based UCB reaches 90% in 12,000/4,500 rounds (leak-aware UCB: 27,000/7,000) vs. 56,000/18,000 for UCB1	Faster: supported; nears the CCE set: weak*
H5	$\Lambda_{\text{BAR}} > \Lambda_{\text{SEATS}}$, and fixed seats need a smaller Q for $\Lambda < 5\%$	CI includes 0, or order reversed	At $\alpha = 0.2$: 8–12× larger, disjoint CIs at every Q; $Q \geq 9$ vs. $Q \geq 40$. If everyone told: 2.5–4.4×; $Q \geq 20$ vs. $Q \geq 40$	Supported

* The CCE gap falls but is not yet small. The model-based learner is a reduced-form version of the one in the review (it estimates the value β of a leaked question, not Δ, and plans alone rather than best-responding in an estimated game of all candidates); the leak-aware learner is only UCB1 with credit for tells that did not land (Section 4) and has lower regret only under a cut-off.

is instead reported per equilibrium by the bootstrap; (ii) the "model-based" learner of the review is built in two simpler forms: Leak-aware UCB (UCB1 with counterfactual credit; its leak-rate estimate is logged but not used) and Model-based UCB (a reduced-form model of leaked questions, planned with by one candidate alone); both need the friend to report how many told questions came up; (iii) mixed equilibria were sought by replicator dynamics, by iterated dominance on all networks and by support enumeration on the pairs network only (agents have four or eight actions on the other two), with the CCE programs bounding what remains; (iv) the CCE programs use a coarser grid (α in steps of 0.05) for two friend groups; (v) the complete network at $n = 6$ was handled by deviation tests, and $n = 8$ was not attempted. **Limits.** Types are coarse and rounds are independent: there is no reputation, and placed candidates do not leave the pool (a stochastic game [14]). Networks are fixed and small. There is no real leak data, so every claim is comparative (seats vs. cut-off, weak vs. strong), not a prediction of how much real candidates share. The bootstrap measures Monte Carlo error, not sensitivity to Δ, σ or K/n, which we held fixed.

6 Contribution

Member	Task assigned in the review	Where it is delivered	Status
Jash Bharat Pasad	Formulation, proofs; utility types (X2); company analysis (X4); writing, integration	Sections 2–3, Corollary 1; `utility.py`, `run_x2.py`, `run_x4.py`; Table 2, Figure 5; report	Fulfilled
Bhavay Goyal	Simulator (E0–E2, networks), payoff tables, verification checks, pipeline	`config.py`, `env.py`, `tables.py`, `checks.py`, unit tests; Table 1	Fulfilled
Chaitanya Bhoite	Level-1 solvers (BNE, mixed equilibria, CCE program, bootstrap ε); X1	`game.py`, `solvers.py`, `support_enum.py`, `analysis.py`, `run_x1.py`, `run_cce.py`, `run_support.py`; Figures 1–3	Fulfilled; support enumeration on the pairs network only
Vishnu Sinha	Level-2 learners (incl. model-based); regret and CCE metrics; X3; statistics	`learners.py`, `learning.py`, `stats.py`, `run_x3.py`; Figure 4, Table 3	Fulfilled; model-based learner in reduced form

Core (verified simulator, X1–X3, H1–H4) and target (X4, H5) deliverables are complete; the stretch target $n = 8$ is not. Every number above is generated from the result files in the repository by `experiments/make_figures.py`.

References

[1] Team Blind. Candidate confessed to memorizing leaked interview questions. https://www.teamblind.com, accessed Sept. 2026.

[2] Glider AI. Leaked interview questions: How candidates get answers. https://glider.ai, accessed Sept. 2026.

[3] E. P. Lazear and S. Rosen. Rank-order tournaments as optimum labor contracts. *J. Polit. Econ.*, 89(5):841–864, 1981.

[4] M. Kräkel. Helping and sabotaging in tournaments. *Int. Game Theory Rev.*, 7(2):211–228, 2005.
[5] Y. Li and V. Conitzer. Game-theoretic question selection for tests. *J. Artif. Intell. Res.*, 59:437–462, 2017.
[6] P. Auer, N. Cesa-Bianchi, and P. Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47:235–256, 2002.
[7] P. Auer, N. Cesa-Bianchi, Y. Freund, and R. E. Schapire. The nonstochastic multiarmed bandit problem. *SIAM J. Comput.*, 32(1):48–77, 2002.
[8] W. R. Thompson. On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25(3/4):285–294, 1933.
[9] S. Hart and A. Mas-Colell. A simple adaptive procedure leading to correlated equilibrium. *Econometrica*, 68(5):1127–1150, 2000.
[10] M. P. Wellman. Methods for empirical game-theoretic analysis. In *Proc. AAAI*, pages 1552–1556, 2006.
[11] Y. Shoham and K. Leyton-Brown. *Multiagent Systems: Algorithmic, Game-Theoretic, and Logical Foundations*. Cambridge University Press, 2008.
[12] E. Fehr and K. M. Schmidt. A theory of fairness, competition, and cooperation. *Q. J. Econ.*, 114(3):817–868, 1999.
[13] S. Holm. A simple sequentially rejective multiple test procedure. *Scand. J. Stat.*, 6(2):65–70, 1979.
[14] L. S. Shapley. Stochastic games. *Proc. Natl. Acad. Sci.*, 39(10):1095–1100, 1953.

A Additional results

Table 5: Median rounds to 90% agreement for all three leak models (Table 3 shows E1 only); "never (x%)": the median run did not reach it, x% of runs did.

	E0		E1		E2	
Learner	Seats	Cut-off	Seats	Cut-off	Seats	Cut-off
ε-greedy	never (47%)	never (13%)	never (20%)	never (50%)	never (13%)	160.5k
UCB1	16.5k	never (3%)	56k	18k	60.5k	12k
Thompson sampling	19.5k	never (13%)	78k	19k	80k	16.5k
EXP3	29k	never (3%)	141k	61k	122k	41k
Leak-aware UCB	16.5k	never (3%)	27k	7k	28.5k	8.5k
Model-based UCB	12k	never (0%)	12k	4.5k	15.5k	4k

Table 6: What a sharing pattern is worth (E1, two friend groups, $\alpha = 0.2$): the baselines of the review and the profile each learner settled on (most-played row of every decision cell; mean over 30 seeds; for Leak-aware UCB the profile of greedy means is identical in every run). Welfare is the sum of the six candidates' expected utilities. Under fixed seats exactly K candidates are hired whatever is shared, so welfare is the same for every pattern up to the message cost and is omitted. Company value is $\mathbb{E}[V]$; "nobody shares" is the leak-free company.

	Fixed seats		Cut-off		
Sharing pattern	Tells (%)	Company value	Tells (%)	Welfare	Company value
Nobody shares (leak-free company)	0	0.657	0	2.799	0.894
Random sharing	50	0.631	50	3.067	0.833
Everybody shares	100	0.614	100	3.317	0.771
Equilibrium	33	0.645	100	3.317	0.771
ε-greedy	43	0.638	87	3.256	0.787
UCB1	38	0.643	99	3.314	0.772
Thompson sampling	38	0.643	99	3.314	0.772
EXP3	41	0.639	94	3.288	0.778
Leak-aware UCB	34	0.645	100	3.317	0.771
Model-based UCB	33	0.645	100	3.317	0.771